



Quantitative Project Risk Analysis with PRA

Paul Govan 
GE Aerospace

Abstract

The **PRA** package for R is the first package to unify a core set of widely adopted quantitative project risk analysis methods in a single, open-source implementation: the Second Moment Method (SMM), Monte Carlo simulation (MCS), earned value management (EVM), contingency and sensitivity analysis, noisy-OR risk inference and updating, sigmoidal learning curves, and design structure matrices (DSM). Beyond this foundational core, **PRA** provides an open-source implementation of probabilistic network methods for project risk, supporting the forward simulation of cost uncertainty and Bayesian conditioning along causal chains. **PRA** also exposes every analytical function as a tool over the Model Context Protocol (MCP) so that large-language-model (LLM) agents can run project risk analyses directly. The package is annotated against the rOpenSci general standards for statistical software and is supported by a continuous integration test suite running on macOS, Windows, and Ubuntu.

Keywords: project risk analysis, Monte Carlo simulation, earned value management, noisy-OR inference, design structure matrix, R.

1. Introduction

Quantitative risk analysis is a foundational discipline in project management, enabling organizations to assess schedule and cost uncertainty, identify high-impact risk drivers, and set evidence-based contingency reserves. This article presents **PRA**, an open-source R package that unifies the core quantitative methods of project risk analysis, schedule and cost uncertainty, earned value management, sensitivity and contingency analysis, learning curves, and dependency structure analysis, under a single, reproducible API.

Despite broad industry adoption of quantitative methods, open-source tooling for integrated project risk analysis in R remains fragmented. Proprietary platforms such as Oracle Primavera Risk Analysis and Palisade @RISK provide integrated workflows but impose licensing costs and limit reproducibility. Within R, no single package unifies Monte Carlo risk analysis,

earned value management, noisy-OR risk inference, sigmoidal learning curves, and dependency structure analysis under a common API. The **mc2d** package (Pouillot and Delignette-Muller 2010) provides general-purpose two-dimensional Monte Carlo simulation but has no project-specific constructs. The **sensitivity** package (Iooss, Da Veiga, Janon, and Pujol 2024) implements variance-based indices for general functions but does not integrate with schedule or cost models. EVM metrics are absent from CRAN entirely. Bayesian packages such as **rstan** (Stan Development Team 2024) and **brms** (Bürkner 2017) operate at a high level of generality and require substantial expertise to apply in project risk contexts. For probabilistic networks specifically, **bnlearn** (Scutari 2010) and **gRain** (Højsgaard 2012) are the established R tools: **bnlearn** provides structure and parameter learning along with graph-surgery interventions through **mutilated()**, and **gRain** performs exact junction-tree inference. Both are more general and more capable inference engines than the sampling-based module of Section 6. However, neither carries a project-risk vocabulary of risks, resources, and tasks, and neither connects to schedule, cost, or earned value; **PRA**’s contribution is that project layer rather than the underlying inference.

PRA is the first package available on CRAN to integrate these project risk analysis methods in a single API. The comparison in Table 1 covers the closest general-purpose Monte Carlo and sensitivity packages and the dominant commercial platform, while a broader survey of project-specific R packages is provided below. Where the existing ecosystem is fragmented across general-purpose packages, none of which are project-specific, **PRA** provides a single reproducible, auditable API for the full workflow, integrating outputs across modules in a way that the fragmented ecosystem does not. For example, MCS results are directly fed into contingency and sensitivity analyses, and noisy-OR risk updates can re-parameterize task distributions for subsequent simulations. The package targets project analysts who work in R and require outputs that are compatible with standard reporting pipelines. It is annotated against the rOpenSci Statistical Software Review (SRR) general standards (rOpenSci Statistical Software Peer Review 2021) and uses minimal hard dependencies: **mc2d** (Pouillot and Delignette-Muller 2010) for triangular distribution families and **minpack.lm** (Elzhov, Mullen, Spiess, and Bolker 2016) for nonlinear least-squares fitting.

PRA also provides an R implementation of probabilistic-network methods for project risk that grounds causal networks in the resource-based view, showing that organizational capabilities and shared resources, not schedule uncertainty alone, are the structural pathways through which risks propagate across project tasks (Govan 2014; Govan and Damnjanovic 2016). This approach is extended with structural network measures that quantify the topology of project risk networks (Govan and Damnjanovic 2020). The probabilistic network module (Section 6) operationalizes these ideas by propagating distributions through the causal chain from risks to the total project risk.

PRA is the first project risk analysis package to expose its analytical methods as tools callable by LLM agents over the Model Context Protocol (MCP). The optional MCP server is activated through **Suggests**-only dependencies; therefore, the analytical core incurs no additional runtime cost.

The following table underscores the primary contribution: no existing R package or commercial platform covers the foundational project risk workflow that **PRA** unifies. It summarizes how **PRA** compares with the most closely related R packages and the dominant commercial platform.

Table 1: Comparison of **PRA** with related R packages and commercial tools. “Partial” = only general-purpose support, requiring custom modeling and not project-specific; “—” = not supported.

Feature	PRA	mc2d	sensitivity	rstan/brms	bnlearn/gRain	@RISK
MCS schedule risk	Yes	Partial	—	—	—	Yes
EVM (11 metrics)	Yes	—	—	—	—	Yes
Noisy-OR updating	Yes	—	—	Yes	Partial	—
Learning curves	Yes	—	—	—	—	—
Design structure matrices	Yes	—	—	—	—	—
Probabilistic networks	Yes	—	—	Partial	Yes	—
MCP tool server	Yes	—	—	—	—	—
Open-source (CRAN)	Yes	Yes	Yes	Yes	Yes	—
SRR general standards	Yes	—	—	—	—	—

A small number of CRAN packages do target project management specifically, but each addresses a single slice of the workflow rather than the integrated whole. The **Project-Management** package (Gonçalves-Dosantos, García-Jurado, and Costa 2022) is the closest in spirit: it implements deterministic and stochastic scheduling, PERT/CPM, critical-path identification, and Monte Carlo simulation over activity durations, and is complementary to the schedule-network extension planned for **PRA** (Section 10), but it does not provide earned value management, noisy-OR risk updating, learning curves, dependency structure matrices, or probabilistic networks. The **plan** package (Kelley 2023) produces Gantt and burn-down charts for project visualization but performs no risk quantification. More generally, decision-analysis packages such as **decisionSupport** (Luedeling, Goehring, Schiffers, Whitney, and Fernandez 2023) apply Monte Carlo simulation to cost–benefit models under uncertainty, but are not organized around project-schedule or earned-value constructs. None of these packages, nor any single package on CRAN, spans the foundational project risk workflow that **PRA** unifies.

The gap is not limited to R: open-source tooling for integrated project risk analysis remains sparse across languages, and no package in another ecosystem matches **PRA**’s combined scope either. In Python, **PyCostTools** (Frank 2023) implements learning-curve and cost-estimating-relationship models similar in spirit to Section 4, but remains a low-adoption project without peer-reviewed documentation. The general-purpose **PyMC** (Salvatier, Wiecki, and Fonnesbeck 2016) supports arbitrary Bayesian models but, unlike the probabilistic-network module of Section 6, has no project-risk-specific vocabulary of risks, resources, and tasks. In Julia, **MCHammer.jl** (Torkia 2023) provides Monte Carlo business-risk modeling with correlated sampling and sensitivity tornado charts in the style of Crystal Ball and @RISK, but has seen minimal maintenance activity in recent years. Design structure matrix tools outside R target software and systems architecture rather than project risk, and no citable academic survey of open-source project-risk software outside R could be located; this survey is therefore drawn directly from the individual tools’ own repositories rather than a secondary source.

The package is installed from CRAN in the standard way:

```
R> install.packages("PRA")
R> library("PRA")
```

Optional dependencies, **ellmer** (Wickham, Cheng, Jacobs, Aden-Buie, and Schloerke 2025), **mcptools** (Couch, Chang, and Gao 2025), and **jsonlite** (Ooms 2014), are only required for the MCP server, `pra_mcp_server()`.

The remainder of this article describes each analytical module in turn, with working code examples drawn from the same R session. Section 2 covers schedule and cost uncertainty methods (SMM, MCS, contingency, sensitivity, and correlation matrix construction). Section 3 describes earned value management. Section 4 describes sigmoidal learning curves. Section 5 covers noisy-OR risk inference and updating. Section 6 introduces the probabilistic network module. Section 7 covers design structure matrices, built from that same network example. Section 8 describes the Model Context Protocol integration. Section 9 presents an integrated case study. Section 10 concludes.

2. Schedule and Cost Uncertainty

Two complementary methods are provided for quantifying schedule and cost uncertainty: the Second Moment Method for rapid analytical estimates, and Monte Carlo simulation for full distributional analysis. Contingency extraction and sensitivity analysis complete the workflow.

To illustrate the workflow, **PRA** includes `building_project`, a bundled example dataset describing a mid-rise commercial building with six work packages, adapted from the construction-project examples in Damnjanovic and Reinschmidt (2020).

```
R> data("building_project", package = "PRA")
R> bp <- building_project
R> bp$task_names
```

```
[1] "Site Preparation & Earthworks" "Foundations"
[3] "Structural Frame"             "Building Envelope"
[5] "Mechanical/Electrical/Plumbing" "Interior Fit-out & Finishes"
```

2.1. Second Moment Method

The Second Moment Method (SMM) propagates task-level means and variances to project-level estimates using the Central Limit Theorem (Benjamin and Cornell 2000). For a project with n tasks:

$$E[X] = \sum_{i=1}^n E[X_i], \quad \text{Var}(X) = \sum_{i=1}^n \text{Var}(X_i) + 2 \sum_{i < j} \rho_{ij} \sigma_i \sigma_j$$

The six task durations are triangular, so their means and variances may be obtained in closed-form and propagated with `smm()` at the project level.

```
R> tri_mean <- function(d) (d$a + d$b + d$c) / 3
R> tri_var <- function(d) (d$a^2 + d$b^2 + d$c^2 -
+                          d$a * d$b - d$a * d$c - d$b * d$c) / 18
R> task_means <- vapply(bp$task_distributions, tri_mean, numeric(1))
```

```
R> task_vars <- vapply(bp$task_distributions, tri_var, numeric(1))
R>
R> result <- smm(task_means, task_vars)
R> cat("Total mean:", round(result$total_mean, 2), "weeks\n")
```

Total mean: 71.67 weeks

```
R> cat("Total SD: ", round(result$total_std, 2), "weeks\n")
```

Total SD: 4.53 weeks

Positive correlation between tasks inflates the total variance; `smm()` accepts a correlation matrix and adds the covariance terms analytically:

```
R> result_cor <- smm(task_means, task_vars, bp$cor_mat)
R> cat("Total SD (correlated):", round(result_cor$total_std, 2), "weeks\n")
```

Total SD (correlated): 6.59 weeks

SMM is useful when only first and second moments are available and a rapid estimate is needed. When task distributions are non-normal or accurate tail behavior is required, Monte Carlo simulation is preferred.

2.2. Monte Carlo Simulation

`mcs()` propagates uncertainty through a project network by drawing samples from user-specified task distributions, triangular, normal, or uniform, over a large number of iterations (Vose 2008). By default the tasks are sampled independently, with an optional correlation matrix that induces dependence between them. For a project with n tasks, one simulated draw of the total duration is

$$T^{(s)} = \sum_{i=1}^n Y_i^{(s)}, \quad s = 1, \dots, S,$$

where $Y_i^{(s)}$ is the s -th draw of task i 's duration, and the project-level estimates are the Monte Carlo moments of the S simulated totals,

$$\hat{E}[T] = \frac{1}{S} \sum_{s=1}^S T^{(s)}, \quad \widehat{\text{Var}}(T) = \frac{1}{S-1} \sum_{s=1}^S (T^{(s)} - \hat{E}[T])^2.$$

When a correlation matrix $\mathbf{C}$ is supplied, dependence between tasks is induced by Cholesky decomposition. The correlation matrix is factored as $\mathbf{U}^\top \mathbf{U} = \mathbf{C}$, and a matrix $\mathbf{Z}$ of independent standard normal draws (columns = tasks) is right-multiplied by the Cholesky factor $\mathbf{U}$, producing normal scores that carry exactly the target correlation. Because task distributions are in general not normal, the correlated scores are then returned to each task's own scale: the standard normal CDF Φ maps each score to a percentile, and each task's inverse CDF F_i^{-1} maps that percentile back to a duration,

$$\mathbf{W} = \mathbf{Z}\mathbf{U}, \quad U_i^{(s)} = \Phi(W_i^{(s)}), \quad Y_i^{(s)} = F_i^{-1}(U_i^{(s)}).$$

```
R> set.seed(42)
R> results <- mcs(10000, bp$task_distributions)
R> cat("Simulated mean:", round(results$total_mean, 2), "weeks\n")
```

Simulated mean: 71.72 weeks

```
R> cat("Simulated SD: ", round(results$total_sd, 2), "weeks\n")
```

Simulated SD: 4.57 weeks

The `mcs()` function returns an S3 object of class "mcs" with a `print()` method that summarizes percentiles. The full simulated distribution is available in `results$total_distribution`. On a modern laptop, 10,000 simulations of a six-task network complete in under one second; increasing the count to 100,000 provides more precise tail estimates at roughly ten-fold additional runtime. The default of 10,000 balances accuracy and speed for most practical analyses. Supplying the dataset's correlation matrix induces dependence between the tasks through the transform above:

```
R> set.seed(42)
R> results_cor <- mcs(10000, bp$task_distributions, bp$cor_mat)
R> cat("Simulated SD (correlated):", round(results_cor$total_sd, 2), "weeks\n")
```

Simulated SD (correlated): 6.57 weeks

Numerical Validation

For independent tasks the SMM estimates are the exact first two moments of the total: the total mean equals the sum of the task means by linearity of expectation, and the total variance equals the sum of the task variances by additivity of variance under independence. Supplying the correlation matrix adds the covariance terms analytically, giving a second exact target. Monte Carlo simulation should reproduce both, and it does, to within Monte Carlo error. The correlated comparison is the sharper test of the two: an error in the Cholesky factorization or in the rank-to-scale mapping would surface here as a mismatched total variance, while leaving the independent case untouched.

```
R> cat(sprintf("Mean:           SMM %.2f  vs  MCS %.2f\n",
+             result$total_mean, results$total_mean))
```

Mean: SMM 71.67 vs MCS 71.72

```
R> cat(sprintf("SD:           SMM %.2f  vs  MCS %.2f\n",
+             result$total_std, results$total_sd))
```

SD: SMM 4.53 vs MCS 4.57




Figure 1: Distribution of simulated total project durations. The dashed vertical line marks the mean.

```
R> cat(sprintf("SD (correlated): SMM %.2f vs MCS %.2f\n",
+             result_cor$total_std, results_cor$total_sd))

SD (correlated): SMM 6.59 vs MCS 6.57

R> hist(results$total_distribution,
+       breaks = 50, freq = FALSE,
+       main = "MCS: Total Project Duration",
+       xlab = "Duration (weeks)", col = "steelblue", border = "white"
+ )
R> lines(density(results$total_distribution), col = "tomato", lwd = 2)
R> abline(v = results$total_mean, lty = 2, lwd = 1.5)
```

2.3. Correlation Matrix Construction

Both `smm()` and `mcs()` accept a correlation matrix to model dependence between tasks. `cor_matrix()` supplies one by sampling from a set of task distributions, giving a correctly shaped, positive-definite baseline for testing; correlations reflecting genuine shared drivers are best drawn from historical data or expert elicitation. For samples x_{ki} , $k = 1, \dots, N$, drawn from each of m task distributions $i = 1, \dots, m$, `cor_matrix()` returns the empirical Pearson correlation matrix $\mathbf{C}$ with entries

$$C_{ij} = \frac{\sum_{k=1}^N (x_{ki} - \bar{x}_i)(x_{kj} - \bar{x}_j)}{\sqrt{\sum_{k=1}^N (x_{ki} - \bar{x}_i)^2} \sqrt{\sum_{k=1}^N (x_{kj} - \bar{x}_j)^2}}.$$

```
R> set.seed(42)
R> dists <- list(
+   normal = function(n) rnorm(n, mean = 10, sd = 2),
+   triang = function(n) mc2d::rpert(n, min = 5, mode = 10, max = 15),
+   uniform = function(n) runif(n, min = 8, max = 12)
+ )
R> C_emp <- cor_matrix(num_samples = 1000, num_vars = 3, dists = dists)
R> print(round(C_emp, 3))
```

```
      normal triang uniform
normal  1.000 -0.038 -0.019
triang -0.038  1.000  0.053
uniform -0.019  0.053  1.000
```

2.4. Contingency Analysis

`contingency()` extracts a reserve buffer as the difference between a high-confidence percentile and the base (median) estimate, a standard approach in project risk practice ([Project Management Institute 2021](#)). Writing $Q_p(T)$ for the empirical p -th quantile of the simulated total duration T , the contingency reserve is

$$\text{CR} = Q_{p_{\text{high}}}(T) - Q_{p_{\text{base}}}(T).$$

```
R> reserve <- contingency(results, phigh = 0.80, pbase = 0.50)
R> cat("Schedule contingency (P80 - P50):", round(reserve, 2), "weeks\n")
```

```
Schedule contingency (P80 - P50): 3.98 weeks
```

2.5. Sensitivity Analysis

`sensitivity()` decomposes total project variance by task contribution, producing the inputs for a Tornado chart. Each task's index is its own variance plus its covariance with every other task, expressed as a proportion of total project variance. Let $\text{Var}(X_i)$ be task i 's variance and $\text{Cov}(X_i, X_j) = \rho_{ij} \sqrt{\text{Var}(X_i) \text{Var}(X_j)}$ its covariance with task j (zero if no correlation matrix is supplied). Total project variance decomposes as

$$\text{Var}(T) = \sum_{i=1}^n \text{Var}(X_i) + 2 \sum_{i < j} \text{Cov}(X_i, X_j),$$

and `sensitivity()` reports each task's share of that total,

$$S_i = \frac{\text{Var}(X_i) + \sum_{j \neq i} \text{Cov}(X_i, X_j)}{\text{Var}(T)}, \quad \sum_{i=1}^n S_i = 1.$$

```
R> sens <- sensitivity(bp$task_distributions)
R> names(sens) <- paste0("T", seq_along(sens))
```





Figure 2: Tornado chart: each bar shows a task's share of total project duration variance, sorted ascending so the largest driver appears at the top.

```
R> barplot(sort(sens),
+   horiz = TRUE, col = "steelblue",
+   main = "Sensitivity: Task Contributions",
+   xlab = "Proportion of Total Variance"
+ )
```

Tasks with larger bars are the dominant drivers of total risk; mitigation resources should be directed there first.

3. Earned Value Management

The EVM module implements the complete ANSI/EIA-748 suite of eleven performance metrics (Fleming and Koppelman 2010). The three core quantities, Planned Value (PV), Earned Value (EV), and Actual Cost (AC), are computed first. All performance indices and forecasts follow from them. Consider a project with a \$500,000 budget at completion, evaluated at period 3 of a 5-period schedule with cumulative planned value 10/25/50/75/100%:

```
R> bac      <- 500000
R> schedule <- c(0.10, 0.25, 0.50, 0.75, 1.00)
R> time_period <- 3
```

At this reporting point, 40% of scope is actually complete, and actual costs of \$45,000, \$110,000, and \$135,000 have been incurred across the three periods to date:

```
R> pv_val <- pv(bac, schedule, time_period)
R> ev_val <- ev(bac, actual_per_complete = 0.40)
R> ac_val <- ac(c(45000, 110000, 135000), time_period, cumulative = FALSE)
R>
R> cat("PV: $", format(pv_val, big.mark = ","), "\n")
```

PV: \$ 250,000

```
R> cat("EV: $", format(ev_val, big.mark = ","), "\n")
```

EV: \$ 200,000

```
R> cat("AC: $", format(ac_val, big.mark = ","), "\n")
```

AC: \$ 290,000

The Schedule Performance Index, Cost Performance Index, Schedule Variance, and Cost Variance follow directly from PV, EV, and AC:

$$SPI = \frac{EV}{PV}, \quad CPI = \frac{EV}{AC}, \quad SV = EV - PV, \quad CV = EV - AC$$

```
R> spi_val <- spi(ev_val, pv_val)
R> cpi_val <- cpi(ev_val, ac_val)
R> sv_val <- sv(ev_val, pv_val)
R> cv_val <- cv(ev_val, ac_val)
R>
R> cat("SPI:", round(spi_val, 3), " CPI:", round(cpi_val, 3), "\n")
```

SPI: 0.8 CPI: 0.69

```
R> cat("SV: $", format(round(sv_val), big.mark = ","),
+      " CV: $", format(round(cv_val), big.mark = ","), "\n")
```

SV: \$ -50,000 CV: \$ -90,000

Both $SPI < 1$ and $CPI < 1$ indicate the project is behind schedule and over budget. The `eac()` function provides three completion forecasts under different efficiency assumptions ([Fleming and Koppelman 2010](#)):

$$EAC_{\text{typical}} = \frac{BAC}{CPI}, \quad EAC_{\text{atypical}} = AC + (BAC - EV), \quad EAC_{\text{combined}} = AC + \frac{BAC - EV}{CPI \times SPI}$$

The typical method assumes remaining work proceeds at the cumulative cost efficiency to date; the atypical method assumes the overrun to date was a one-off and future work runs on plan; the combined method deflates the remaining budget by both cost and schedule performance.

```
R> cat("EAC (typical): $", format(round(eac(bac, method = "typical",
+   cpi = cpi_val))), big.mark = ","), "\n")
```

```
EAC (typical): $ 725,000
```

```
R> cat("EAC (atypical): $", format(round(eac(bac, method = "atypical",
+   ac = ac_val, ev = ev_val))), big.mark = ","), "\n")
```

```
EAC (atypical): $ 590,000
```

```
R> cat("EAC (combined): $", format(round(eac(bac, method = "combined",
+   cpi = cpi_val, ac = ac_val, ev = ev_val, spi = spi_val))),
+   big.mark = ","), "\n")
```

```
EAC (combined): $ 833,750
```

The remaining EVM metrics are:

```
R> eac_typ <- eac(bac, method = "typical", cpi = cpi_val)
R> cat("ETC: $",
+   format(round(etc(bac, ev_val, cpi = cpi_val))), big.mark = ","), "\n")
```

```
ETC: $ 435,000
```

```
R> cat("TCPI:", round(tcpi(bac, ev_val, ac_val), 3), "\n")
```

```
TCPI: 1.429
```

```
R> cat("VAC: $", format(round(vac(bac, eac_typ))), big.mark = ","), "\n")
```

```
VAC: $ -225,000
```

The To-Complete Performance Index (TCPI) reports the cost efficiency required on all remaining work to finish within the original budget; values substantially above 1.0 signal an unrealistic recovery target. The Estimate to Complete (ETC) gives the projected cost of the remaining work under the typical-method efficiency assumption, and the negative Variance at Completion (VAC) confirms the project is tracking toward a budget overrun rather than an underrun.

4. Learning Curves

Sigmoidal (S-curve) models capture the characteristic pattern of project progress: slow start, rapid acceleration, and plateau approaching 100% completion. **PRA** provides two model types, Logistic and Gompertz, fitted via the Levenberg–Marquardt nonlinear least-squares algorithm in **minpack.lm** (Elzhov *et al.* 2016). The Logistic model has the closed form

$$y(t) = \frac{K}{1 + \exp(-r(t - t_0))},$$

while the Gompertz model is

$$y(t) = A \exp(-b \exp(-ct)).$$

Each model is fit by nonlinear least squares,

$$\hat{\theta} = \arg \min_{\theta} \sum_{k=1}^m (y_k - f(t_k; \theta))^2,$$

where θ is (K, r, t_0) or (A, b, c) and f is the corresponding curve.

```
R> progress <- data.frame(
+   time      = 1:9,
+   completion = c(5, 15, 40, 60, 70, 75, 80, 85, 90)
+ )

R> fit_log <- fit_sigmodal(progress, "time", "completion", "logistic")
R> fit_gomp <- fit_sigmodal(progress, "time", "completion", "gompertz")
R>
R> cat("Logistic RSE: ", round(summary(fit_log)$sigma, 2), "\n")
```

Logistic RSE: 4.15

```
R> cat("Gompertz RSE: ", round(summary(fit_gomp)$sigma, 2), "\n")
```

Gompertz RSE: 2.89

The Gompertz model has the lower residual standard error of the two here, reflecting the asymmetric approach to the 90% plateau in the observed data, and is therefore carried forward for the confidence band and forecasts below. The Logistic form remains available via `fit_sigmodal()`. Its symmetric acceleration and deceleration is the more common default assumption for project progress curves, and comparing the two residual standard errors is the recommended way to choose between them in practice.

```
R> plot_sigmodal(fit_gomp, progress, "time", "completion", "gompertz",
+               conf_level = 0.95,
+               main = "Gompertz Learning Curve",
+               xlab = "Week", ylab = "Completion (%)")
```

`predict_sigmodal()` generates forecasts with confidence bounds at arbitrary future time points:

```
R> preds <- predict_sigmodal(fit_gomp, seq(10, 12, by = 0.5),
+                           "gompertz", conf_level = 0.95)
R> print(round(preds, 1))
```




Figure 3: Gompertz learning curve over the observed data range, with a 95% confidence band. The band is narrowest in the middle of the range, where the observations most tightly constrain the fit, and widens toward both ends.

	x	pred	lwr	upr
1	10.0	86.8	81.9	91.7
2	10.5	87.0	81.9	92.0
3	11.0	87.1	81.9	92.3
4	11.5	87.2	81.9	92.5
5	12.0	87.3	81.9	92.6

The forecast extends the fitted curve past the last observed week (9) out to week 12, with intervals of the form $\text{pred} \pm t\text{se}$ that widen with the horizon. The point estimates rise only slightly across that span, since the curve is already near its fitted asymptote $\hat{A} \approx 87$, which on nine observations is best read as a lower bound on eventual completion.

5. Noisy-OR Risk Inference

Project risks are frequently triggered by one or more underlying root causes whose presence is uncertain. **PRA** models this with a noisy-OR risk network (Pearl 1988): each root cause C_i acts on a risk event R independently, and R occurs if it is triggered by at least one cause. Write q_i^+ for the probability that cause i alone triggers R when it is present and q_i^- for the corresponding probability when it is absent; these are the link probabilities supplied as `risks_given_causes` and `risks_given_not_causes`. Note that q_i^+ is not the model's conditional $P(R \mid C_i)$, which is larger because the remaining causes may also fire. We illustrate with the schedule-delay risk in `building_project`, which has two candidate causes.

5.1. Prior Risk Probability

`risk_prob()` first computes each cause's marginal contribution by the law of total probability,

$$m_i = q_i^+ P(C_i) + q_i^- (1 - P(C_i)),$$

and then combines the independent causes with a noisy-OR,

$$P(R) = 1 - \prod_{i=1}^k (1 - m_i),$$

which always lies in $[0, 1]$ and reduces to the single-cause marginal when $k = 1$.

```
R> bp <- building_project
R> prior_delay <- risk_prob(bp$cause_probs, bp$risks_given_causes,
+                           bp$risks_given_not_causes)
R> cat("Prior P(schedule delay):", round(prior_delay, 3), "\n")
```

```
Prior P(schedule delay): 0.578
```

5.2. Posterior Risk Probability

When some causes are subsequently observed, for example, a site inspection confirms one cause is present while the other remains unknown, `risk_post_prob()` updates the estimate. An observed cause is fixed to its conditional probability, while an unobserved cause (coded NA) retains its marginal contribution m_i ; the causes are recombined with the same noisy-OR, so prior and posterior share a common scale:

$$P(R \mid \text{obs}) = 1 - \prod_{i=1}^k (1 - c_i), \quad c_i = \begin{cases} q_i^+ & C_i \text{ observed present,} \\ q_i^- & C_i \text{ observed absent,} \\ m_i & C_i \text{ unobserved.} \end{cases}$$

```
R> post_delay <- risk_post_prob(
+   bp$cause_probs, bp$risks_given_causes,
+   bp$risks_given_not_causes, bp$observed_causes
+ )
R> cat("Posterior P(schedule delay):", round(post_delay, 3),
+     " (observed causes:", paste(bp$observed_causes, collapse = ", "), ")\n")
```

```
Posterior P(schedule delay): 0.797 (observed causes: 1, NA )
```

Observing an aggravating cause raises the posterior risk probability; observing a mitigating one lowers it. The update conditions on the observed causes within a fixed model rather than revising beliefs about the model's own parameters, and the diagnostic direction, $P(C_i \mid R)$, is not computed. This closed-form update is inexpensive and complements the sample-based conditioning of the probabilistic network module (Section 6), which operates on full cost distributions rather than event probabilities.

6. Probabilistic Networks

This module is an implementation of probabilistic-network methods for project risk (Govan 2014; Govan and Damnjanovic 2016, 2020). The module is not a general Bayesian-network engine, since other packages already fill that role in R **bnlearn** and **gRain**. What follows is the project-risk layer, in which nodes are risks, resources, tasks, and the project total, and in which the quantity propagated is cost. The probabilistic network is a directed acyclic graph in which nodes represent project variables (risks, resources, tasks, and the project total) and edges encode conditional dependencies (Pearl 2009), attaching a probability distribution to each node so that cost uncertainty can be propagated forward along the causal chain from risks to total project cost and conditioned on observed evidence. Section 7 revisits this same example network from a structural perspective, asking which tasks share resources or risk exposure rather than how their costs propagate.

6.1. Building a Network

`prob_net()` constructs the network from a node data frame, a links data frame, and a named list of distributions. Distributions may be **normal**, **lognormal**, **uniform**, **discrete**, or **conditional** (switching between two sub-distributions depending on a discrete parent node); **aggregate** nodes sum their parents. A probabilistic network encodes a joint distribution over nodes $X_1, \dots, X_n$ that factorizes according to the DAG's parent structure,

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i \mid \text{pa}(x_i)),$$

where $\text{pa}(x_i)$ denotes the values of X_i 's parent nodes (the **condition** node for a conditional distribution, or the summed parents for an aggregate node) (Pearl 2009). The links are load-bearing rather than descriptive: `prob_net()` requires the edges to match the dependencies the distributions declare, so a graph that disagrees with its distributions is rejected at construction, and it requires the nodes to be supplied in a topological order, which is the order `prob_net_sim()` samples in and which also guarantees the graph is acyclic.

The example network below has nine nodes: two risks (Risk-1, Risk-2), three resources (Resource-1, Resource-2, Resource-3), three tasks (Task-1, Task-2, Task-3), and the project total. Risk-1 occurs with probability 0.70 and Risk-2 with probability 0.60; each risk feeds one resource, whose cost distribution switches to a higher-mean, higher-variance branch when its risk occurs, while Resource-3 carries no risk dependency. Each resource aggregates one-to-one into a task (Task-1 from Resource-1, Task-2 from Resource-2, Task-3 from Resource-3), and the three tasks sum into the project total.

```
R> nodes <- data.frame(
+   id = c("A", "B", "C", "D", "E", "F", "G", "H", "I"),
+   label = c("Risk-1", "Risk-2", "Resource-1", "Resource-2", "Resource-3",
+             "Task-1", "Task-2", "Task-3", "Project"),
+   stringsAsFactors = FALSE
+ )
R> links <- data.frame(
+   source = c("A", "B", "C", "D", "E", "F", "G", "H"),
```

```

+   target = c("C","D","F","G","H","I","I","I"),
+   stringsAsFactors = FALSE
+ )
R> distributions <- list(
+   A = list(type = "discrete", values = c(1, 0), probs = c(0.70, 0.30)),
+   B = list(type = "discrete", values = c(1, 0), probs = c(0.60, 0.40)),
+   C = list(type = "conditional", condition = "A",
+     true_dist = list(type = "normal", mean = 30000, sd = 8000),
+     false_dist = list(type = "normal", mean = 15000, sd = 3000)),
+   D = list(type = "conditional", condition = "B",
+     true_dist = list(type = "normal", mean = 80000, sd = 20000),
+     false_dist = list(type = "normal", mean = 50000, sd = 10000)),
+   E = list(type = "normal", mean = 20000, sd = 4000),
+   F = list(type = "aggregate", nodes = "C"),
+   G = list(type = "aggregate", nodes = "D"),
+   H = list(type = "aggregate", nodes = "E"),
+   I = list(type = "aggregate", nodes = c("F","G","H"))
+ )
R> net <- prob_net(nodes, links, distributions = distributions)

```

6.2. Forward Simulation

`prob_net_sim()` draws samples from every node in the order the nodes are supplied, which must be a topological order of the graph, and returns a data frame with one column per node.

```

R> set.seed(42)
R> sim_net <- prob_net_sim(net, num_samples = 10000)
R> cat("Mean project cost: $",
+     format(round(mean(sim_net$I)), big.mark = ","), "\n")

```

Mean project cost: \$ 113,382

```

R> cat("P80 project cost: $", format(round(quantile(sim_net$I, 0.80)),
+     big.mark = ","), "\n")

```

P80 project cost: \$ 135,311

The gap between the mean and the P80 project cost reflects the upper tail introduced by the two risks (nodes A and B): when a risk occurs, its downstream resource is drawn from the higher-cost conditional branch, pulling the tail of the project total upward relative to the mean.

6.3. Learning: Conditioning on Evidence

`prob_net_learn()` draws whole joint samples from the prior and retains only those consistent with the observed values, an observational update in the sense of [Pearl \(2009\)](#). Given M

prior draws $x^{(1)}, \dots, x^{(M)} \sim p(x)$ and observed nodes $X_E = e$, it estimates the conditional expectation of any function g of the network by rejection sampling,

$$\hat{E}[g(X) \mid X_E = e] = \frac{\sum_{m=1}^M \mathbb{I}\{x_E^{(m)} = e\} g(x^{(m)})}{\sum_{m=1}^M \mathbb{I}\{x_E^{(m)} = e\}}.$$

The following call conditions on Risk-2 (node B) being absent (value 0), ruling out Technical Complexity as a concern:

```
R> learn_net <- prob_net_learn(net, observations = list(B = 0),
+                               num_samples = 10000)
R>
R> prior_net_dens <- density(sim_net$I)
R> post_net_dens <- density(learn_net$I)
R>
R> plot(prior_net_dens, col = "steelblue", lwd = 2,
+       main = "Total Project Cost: Prior vs. Posterior",
+       xlab = "Cost ($)",
+       xlim = range(c(sim_net$I, learn_net$I)),
+       ylim = range(c(prior_net_dens$y, post_net_dens$y)))
R> lines(post_net_dens, col = "tomato", lwd = 2)
R> legend("topright",
+       legend = c("Prior (Risk-2 uncertain)", "Posterior (Risk-2 = 0)"),
+       col = c("steelblue", "tomato"), lwd = 2, bty = "n")
```

6.4. Intervention: The do-Operator

Conditioning answers “what if we *observe* that a risk did not occur?” A distinct question is “what if we *act* to eliminate it?” `prob_net_update()` performs structural surgery on the network, adding or removing edges and replacing node distributions, which corresponds to the *do*-operator in Pearl’s causal calculus (Pearl 2009), as opposed to `prob_net_learn()`, which conditions without altering the graph. Intervening to fix $X_k = x_k^*$ replaces X_k ’s conditional distribution with a point mass at x_k^* and removes its incoming edges, giving Pearl’s truncated factorization,

$$p_{\text{do}(X_k=x_k^*)}(x_1, \dots, x_n) = \mathbb{I}\{x_k = x_k^*\} \prod_{i \neq k} p(x_i \mid \text{pa}(x_i)).$$

Suppose management funds a mitigation that decouples Resource-2 from Risk-2: remove the Risk-2 $\rightarrow$ Resource-2 edge and fix Resource-2 to its baseline cost.

```
R> inter_net <- prob_net_update(net,
+   remove_links = data.frame(source = "B", target = "D"),
+   update_distributions = list(
+     D = list(type = "normal", mean = 50000, sd = 10000)
+   )
+ )
R> do_net <- prob_net_sim(inter_net, num_samples = 10000)
```




Figure 4: Total project cost before and after observing that Risk-2 (Technical Complexity) did not occur. Conditioning on Risk-2 = 0 eliminates the heavy upper tail and shifts the distribution leftward by roughly \$18,000 in expectation.

```
R>
R> cat("Prior mean:          $",
+      format(round(mean(sim_net$I)), big.mark = ","), "\n")

Prior mean:          $ 113,382

R> cat("Seeing (Risk-2 = 0):  $",
+      format(round(mean(learn_net$I)), big.mark = ","), "\n")

Seeing (Risk-2 = 0):  $ 95,487

R> cat("Doing (do(Risk-2=0)): $",
+      format(round(mean(do_net$I)), big.mark = ","), "\n")

Doing (do(Risk-2=0)): $ 95,392
```

Here *seeing* (`prob_net_learn()`) and *doing* (`prob_net_update()`) coincide because Risk-2 is an independent root cause with no shared upstream confounder, so conditioning on it and intervening on it induce the same downstream distribution. The two diverge, however, when a risk shares a common cause with project cost, as the next example shows.

6.5. When Seeing and Doing Diverge

Suppose a latent common cause, adverse site conditions, drives both a rework risk and, through a separate pathway, the foundation cost. Observing that rework did *not* occur is evidence that site conditions were favorable, which in turn lowers the expected foundation cost; *intervening* to prevent rework severs the risk but leaves site conditions, and hence the foundation cost, untouched. The network below encodes this structure: the common cause U feeds both the rework risk R_k and the foundation cost W_f , while the schedule cost V_s depends only on R_k .

```
R> set.seed(42)
R> cf_nodes <- data.frame(
+   id = c("U", "Rk", "Wf", "Vs", "Tc"), stringsAsFactors = FALSE)
R> cf_links <- data.frame(
+   source = c("U", "U", "Rk", "Wf", "Vs"),
+   target = c("Rk", "Wf", "Vs", "Tc", "Tc"), stringsAsFactors = FALSE)
R> cf_dists <- list(
+   U = list(type = "discrete", values = c(1, 0), probs = c(0.5, 0.5)),
+   Rk = list(type = "conditional", condition = "U",
+     true_dist = list(type = "discrete",
+       values = c(1, 0), probs = c(0.9, 0.1)),
+     false_dist = list(type = "discrete",
+       values = c(1, 0), probs = c(0.1, 0.9))),
+   Wf = list(type = "conditional", condition = "U",
+     true_dist = list(type = "normal", mean = 60000, sd = 5000),
+     false_dist = list(type = "normal", mean = 20000, sd = 5000)),
+   Vs = list(type = "conditional", condition = "Rk",
+     true_dist = list(type = "normal", mean = 40000, sd = 5000),
+     false_dist = list(type = "normal", mean = 10000, sd = 5000)),
+   Tc = list(type = "aggregate", nodes = c("Wf", "Vs")))
R> cf_net <- prob_net(cf_nodes, cf_links, distributions = cf_dists)
R>
R> cf_prior <- prob_net_sim(cf_net, num_samples = 10000)
R> cf_see <- prob_net_learn(cf_net, observations = list(Rk = 0),
+   num_samples = 10000)
R> cf_do <- prob_net_sim(prob_net_update(cf_net,
+   remove_links = data.frame(source = "U", target = "Rk"),
+   update_distributions = list(
+     Rk = list(type = "discrete", values = c(1, 0), probs = c(0, 1)))),
+   num_samples = 10000)
R>
R> cat("Prior mean total:      $",
+   format(round(mean(cf_prior$Tc)), big.mark = ","), "\n")

Prior mean total:      $ 65,003

R> cat("Seeing (Rk = 0) mean:    $",
+   format(round(mean(cf_see$Tc)), big.mark = ","),
+   " E[U | Rk=0] =", round(mean(cf_see$U), 2), "\n")
```

Seeing (Rk = 0) mean: \$ 33,936 $E[U \mid Rk=0] = 0.1$

```
R> cat("Doing do(Rk = 0) mean: $",
+      format(round(mean(cf_do$Tc)), big.mark = ","),
+      " E[U] =", round(mean(cf_do$U), 2), "\n")
```

Doing do(Rk = 0) mean: \$ 49,798 $E[U] = 0.49$

Conditioning propagates the evidence back through the shared cause, the expected site-condition indicator falls from its prior of 0.5 to about 0.1, so *seeing* that rework did not occur pulls the foundation cost down as well and yields a markedly lower expected total than *doing*. Because `prob_net_learn()` conditions by rejection sampling over whole joint draws, it captures this backdoor propagation that a clamp-and-forward update cannot; this is precisely where observation and intervention must be distinguished, and where the *do*-operator becomes indispensable.

6.6. Ranking Risks by Intervention

Repeating the intervention for each root risk and measuring the reduction in project-cost variance yields a principled, network-level mitigation ranking, a causal analog of the Tornado chart of Section 2:

```
R> base_var <- var(sim_net$I)
R>
R> do_A <- prob_net_sim(prob_net_update(net,
+   remove_links = data.frame(source = "A", target = "C"),
+   update_distributions = list(
+     C = list(type = "normal", mean = 15000, sd = 3000))),
+   num_samples = 10000)
R>
R> importance <- c(
+   "Risk-1" = base_var - var(do_A$I),
+   "Risk-2" = base_var - var(do_net$I)
+ )
R> barplot(sort(importance), horiz = TRUE, col = "steelblue",
+   main = "Risk Importance: Variance Eliminated by do(risk = 0)",
+   xlab = "Reduction in Project-Cost Variance")
```

The probabilistic network module is newer than the analytical core, and its API may still evolve in future releases.

7. Dependency Structure Analysis

Building on foundational DSM literature (Steward 1981; Browning 2001), the DSM module quantifies structural coupling between project tasks through shared resources and risks. This section revisits the probabilistic network of Section 6, with its two risks (Risk-1, Risk-2),

Risk Importance: Variance Eliminated by do(risk = 0)




Figure 5: Risk importance for the probabilistic network: the reduction in total project-cost variance achieved by intervening to eliminate each root risk. Taller bars indicate higher-priority mitigation targets.

three resources (Resource-1, Resource-2, Resource-3), and three tasks (Task-1, Task-2, Task-3), from a structural rather than causal perspective. The network's cost-aggregation edges route each resource's cost to exactly one task for simulation simplicity. In practice, resources and risk exposure are shared more broadly, and the assignment below captures that broader structural coupling.

7.1. Design Structure Matrices

The **Parent DSM** is derived from the Resource-Task matrix $\mathbf{S}$ (rows = resources, columns = tasks) as $\mathbf{P} = \mathbf{S}^\top \mathbf{S}$. Off-diagonal entry P_{jk} counts the number of resources shared between tasks j and k , the structural pathway through which disruptions propagate (Govan and Damnjanovic 2016).

```
R> S <- matrix(c(
+   1, 0, 0,
+   0, 1, 0,
+   1, 1, 1
+ ), nrow = 3, ncol = 3, byrow = TRUE)
R> rownames(S) <- c("Resource-1", "Resource-2", "Resource-3")
R> colnames(S) <- c("Task-1", "Task-2", "Task-3")

R> p <- parent_dsm(S)
R> plot(p)
```




Figure 6: Parent DSM heatmap for the probabilistic network example above. Every task pair shares Resource-3, the common structural bottleneck.

The **Grandparent DSM** adds a risk layer via the Risk–Resource matrix $\mathbf{R}$ (rows = risks, columns = resources), tracing dependencies from risks through resources to tasks as $\mathbf{G} = (\mathbf{RS})^\top (\mathbf{RS})$. Off-diagonal entry $G_{jk} = \sum_i (\mathbf{RS})_{ij} (\mathbf{RS})_{ik}$ scores the risk exposure that tasks j and k share through the resource chain, where $(\mathbf{RS})_{ij}$ counts the risk-to-task paths from risk i to task j . Entries therefore weight a shared risk by the number of common resources carrying it, and reduce to a plain count of shared risks only when $\mathbf{RS}$ is binary. This is a direct operationalization of the structural network measures developed by Govan and Damnjanovic (2020).

```
R> R_mat <- matrix(c(
+   1, 0, 1,
+   0, 1, 1
+ ), nrow = 2, ncol = 3, byrow = TRUE)
R> rownames(R_mat) <- c("Risk-1", "Risk-2")
R> colnames(R_mat) <- c("Resource-1", "Resource-2", "Resource-3")
```

```
R> g <- grandparent_dsm(S, R_mat)
R> plot(g)
```

Resource-3 is the dominant coupling driver: every task pair shares it, and both Risk-1 and Risk-2 propagate through it, concentrating risk exposure across the whole project. It is therefore the prime candidate for contingency buffering or, if feasible, decoupling by assigning a dedicated resource to at least one of the three tasks.




Figure 7: Grandparent DSM heatmap for the probabilistic network example above. Task pairs inherit shared risk exposure through Resource-3.

8. Model Context Protocol Integration

PRA exposes all analytical functions as tools callable by large-language-model (LLM) agents through the Model Context Protocol (MCP) ([Anthropic 2024](#)). `pra_mcp_server()` starts a server that advertises each tool, Monte Carlo simulation, EVM, noisy-OR risk, learning curves, DSM, and the rest, so that MCP-compatible clients can invoke them as native tools. The server is registered in Claude Code with a single shell command:

```
claude mcp add -s project pra -- Rscript -e "PRA::pra_mcp_server()"
```

For Claude Desktop, add the following block to the MCP configuration file:

```
{
  "mcpServers": {
    "pra": { "command": "Rscript",
             "args": ["-e", "PRA::pra_mcp_server()"] }
  }
}
```

The tool definitions are built once with `ellmer::tool()` ([Wickham et al. 2025](#)) and served by **mcptools**. Each tool validates its JSON inputs and returns a plain-text summary for the LLM together with, where applicable, a rich HTML result with inline plots for clients that render it.

The MCP layer is strictly optional: every analytical function operates without any optional dependency. Only `pra_mcp_server()` requires **ellmer**, **mcptools**, and **jsonlite**.

9. Illustrative Case Study

This section demonstrates an integrated, end-to-end project risk analysis for the `building_project` dataset introduced in Section 2, progressing through uncertainty quantification, risk-driver ranking, performance monitoring, forecasting cost at completion, noisy-OR risk updating, dependency structure analysis, and network cost propagation in a single R session. The workflow reuses objects built in the preceding sections (for example the closed-form task moments and the probabilistic network), illustrating how the modules compose within one analysis.

9.1. Step 1: Characterize Task Uncertainty

Combine the rapid analytical estimate from the Second Moment Method, reusing the closed-form triangular moments `task_means` and `task_vars` from Section 2, with the full simulated distribution from Monte Carlo simulation. Both are given the dataset correlation matrix, so the analytical and simulated views rest on the same dependence assumption:

```
R> bp <- building_project
R> set.seed(42)
R>
R> smm_cs <- smm(task_means, task_vars, bp$cor_mat)
R> sim      <- mcs(10000, bp$task_distributions, bp$cor_mat)
R> res      <- contingency(sim, phigh = 0.80, pbase = 0.50)
R>
R> cat("SMM total (correlated): mean", round(smm_cs$total_mean, 1),
+      "SD", round(smm_cs$total_std, 2), "weeks\n")
```

SMM total (correlated): mean 71.7 SD 6.59 weeks

```
R> cat("P50 schedule:",
+      round(quantile(sim$total_distribution, 0.50), 1), "weeks\n")
```

P50 schedule: 71.4 weeks

```
R> cat("P80 schedule:",
+      round(quantile(sim$total_distribution, 0.80), 1), "weeks\n")
```

P80 schedule: 77.3 weeks

```
R> cat("Contingency (P80-P50):", round(res, 2), "weeks\n")
```

Contingency (P80-P50): 5.82 weeks

The P50 estimate is the schedule a planner would commit to with even odds of meeting it; the contingency buffer above it is the additional duration needed to raise that confidence to 80%, and is the reserve a project manager would add to the baseline schedule. Correlation is important in practical terms: treating the six work packages as independent shrinks the same reserve to 3.98 weeks, because independent overruns partially cancel. Shared crews, weather,

and a common supply chain make that assumption optimistic, and the correlated reserve is the defensible one.

9.2. Step 2: Rank Risk Drivers

`sensitivity()` ranks the tasks by their share of total schedule variance, focusing mitigation on the dominant drivers:

```
R> drivers <- sensitivity(bp$task_distributions)
R> names(drivers) <- bp$task_names
R> print(round(sort(drivers, decreasing = TRUE), 3))
```

Structural Frame	Mechanical/Electrical/Plumbing
0.304	0.206
Interior Fit-out & Finishes	Building Envelope
0.206	0.133
Site Preparation & Earthworks	Foundations
0.076	0.076

The Structural Frame accounts for the largest single share of total schedule variance, but it does not dominate: together with the two tasks ranked immediately behind it, it accounts for roughly 70% of the total, and mitigation effort is best spread across all three. The two smallest contributors, at about 8% each, are the least productive targets.

9.3. Step 3: Monitor Earned Value at Mid-Project

Compute planned value, earned value, and actual cost to date for the case-study project:

```
R> pv_cs <- pv(bp$bac, bp$schedule, bp$time_period)
R> ev_cs <- ev(bp$bac, bp$actual_per_complete)
R> ac_cs <- ac(bp$actual_costs, bp$time_period, cumulative = FALSE)
R>
R> cat("SPI:", round(spi(ev_cs, pv_cs), 3),
+      " CPI:", round(cpi(ev_cs, ac_cs), 3), "\n")

SPI: 0.889  CPI: 0.986

R> cat("EAC (typical): $", format(round(eac(bp$bac,
+      method = "typical", cpi = cpi(ev_cs, ac_cs))), big.mark = ","), "\n")

EAC (typical): $ 8,625,000
```

An SPI and CPI both below 1.0 indicate the case-study project is simultaneously behind schedule and over budget at this reporting point, and the typical-method EAC translates that cost inefficiency into a revised completion-cost forecast above the original budget at completion.

9.4. Step 4: Forecast Cost at Completion

The performance indices above translate into a completion forecast. The Estimate to Complete prices the remaining work at the cost efficiency achieved so far, the To-Complete Performance Index gives the efficiency the remaining work would have to achieve to finish within the original budget, and the Variance at Completion states the expected shortfall:

```
R> cpi_cs <- cpi(ev_cs, ac_cs)
R> eac_cs <- eac(bp$bac, method = "typical", cpi = cpi_cs)
R>
R> cat("ETC:  $", format(round(etcc(bp$bac, ev_cs, cpi = cpi_cs)),
+   big.mark = ","), "\n")

ETC:  $ 5,175,000

R> cat("TCPI:", round(tcpi(bp$bac, ev_cs, ac_cs), 3), "\n")

TCPI: 1.01

R> cat("VAC:  $", format(round(vac(bp$bac, eac_cs)), big.mark = ","), "\n")

VAC:  $ -125,000
```

The TCPI is only slightly above 1.0, so the budget is still recoverable: the remaining work has to run about one percent more efficiently than planned, which is a realistic target rather than the kind of step change that signals an unrecoverable overrun. Left uncorrected, the cost efficiency achieved so far projects the modest overrun reported by the Variance at Completion.

9.5. Step 5: Update Risk Probability After Site Inspection

Recompute the noisy-OR prior probability of schedule delay, then update it on the site inspection observation:

```
R> prior_r <- risk_prob(bp$cause_probs, bp$risks_given_causes,
+   bp$risks_given_not_causes)
R> post_r <- risk_post_prob(bp$cause_probs, bp$risks_given_causes,
+   bp$risks_given_not_causes, bp$observed_causes)
R> cat("Prior P(delay):", round(prior_r, 3),
+   "-> Posterior P(delay):", round(post_r, 3), "\n")

Prior P(delay): 0.578 -> Posterior P(delay): 0.797
```

The shift from prior to posterior quantifies how much the site inspection evidence should move the project team's confidence in a schedule delay, information that feeds directly into whether the contingency reserve computed in Step 1 needs to be revisited.

9.6. Step 6: Propagate Cost Through the Network

Use the probabilistic network `net` from Section 6 to propagate cost uncertainty to the project total, then condition on the observation that Risk-2 (Technical Complexity) did not occur:

```
R> cat("Prior mean project cost: $",
+      format(round(mean(sim_net$I)), big.mark = ","), "\n")
```

```
Prior mean project cost: $ 113,382
```

```
R> cat("Prior P80 project cost: $",
+      format(round(quantile(sim_net$I, 0.80)), big.mark = ","), "\n")
```

```
Prior P80 project cost: $ 135,311
```

```
R> cat("Posterior mean (Risk-2 = 0): $",
+      format(round(mean(learn_net$I)), big.mark = ","), "\n")
```

```
Posterior mean (Risk-2 = 0): $ 95,487
```

Conditioning on the resolved risk lowers the expected project cost, closing the loop from uncertainty characterization to evidence-based updating.

9.7. Step 7: Identify Coupled Tasks via DSM

```
R> g_cs <- grandparent_dsm(bp$resource_task, bp$risk_resource)
R> plot(g_cs)
```

The Grandparent DSM reveals which task pairs share risk exposure through common resources, informing targeted contingency allocation and resource decoupling decisions.

10. Summary

PRA is an R package whose primary contribution is being the first package to unify the foundational methods of quantitative project risk analysis in a single, open-source framework of ten integrated analytical modules: SMM, MCS, contingency analysis, sensitivity analysis, EVM (11 metrics), noisy-OR risk inference and updating, cost risk distributions, sigmoidal learning curves, correlation matrix construction, and design structure matrices. The eleventh module applies probabilistic network methods to project risk, with forward simulation and Bayesian conditioning over a network of risks, resources, and tasks.

PRA also exposes every analytical function as a tool over the Model Context Protocol, so that LLM agents can drive the package without any embedded model, the first project risk analysis package to do so. All exported functions are annotated against the rOpenSci SRR general standards, with a test suite verifying correctness, edge-case handling, and numerical stability across three operating systems. The package is licensed under MIT and the source is available at <https://github.com/paulgovan/PRA>; documentation and a companion online book are at <https://paulgovan.github.io/PRA/>.

10.1. Limitations and Future Work




Figure 8: Grandparent DSM for the case-study project, showing which task pairs share the most risk exposure through common resources.

Several limitations apply. First, **PRA** currently propagates uncertainty at the task level and does not yet model schedule logic, critical path, activity precedence, and resource leveling. Full schedule-network risk analysis, in which schedule Monte Carlo runs over an activity-precedence network, such as Program Evaluation and Review Technique / Critical Path Method (PERT/CPM) to identify the critical path and criticality indices under uncertainty, is a planned future direction that would extend **PRA** from task-level to network-level schedule risk. Second, `mcs()` induces task correlation through the Cholesky decomposition of the correlation matrix, which preserves every marginal exactly, but which fixes the dependence structure to the Gaussian family: tail dependence stronger than a multivariate normal implies, such as tasks that fail together only in extreme scenarios, cannot be represented. Richer dependence structures are planned. Third, the probabilistic network module is newer than the analytical core: its API may still evolve. Fourth, the package currently accepts data as R vectors and lists. Direct import from project management file formats (Microsoft Project, Primavera P6, CSV exports from Jira or Asana) is not supported and is planned for a future release.

Computational details

The results in this article were produced using R (R Core Team 2024) with the **PRA** package (version 0.5.0) and its hard dependencies **mc2d** (Pouillot and Delignette-Muller 2010) and **minpack.lm** (Elzhov *et al.* 2016). The optional Model Context Protocol server additionally uses **ellmer** (Wickham *et al.* 2025), **mcptools**, and **jsonlite**. R itself and all packages used are available from the CRAN at <https://CRAN.R-project.org/>. A random seed of 42 is set at

the start of the session so that the simulation output is reproducible.

References

- Anthropic (2024). “Model Context Protocol.” URL <https://modelcontextprotocol.io>.
- Benjamin JR, Cornell CA (2000). *Probability, Statistics, and Decision for Civil Engineers*. Dover Publications, Mineola, NY.
- Browning TR (2001). “Applying the Design Structure Matrix to System Decomposition and Integration Problems: A Review and New Directions.” *IEEE Transactions on Engineering Management*, **48**(3), 292–306. doi:10.1109/17.946528.
- Bürkner PC (2017). “**brms**: An R Package for Bayesian Multilevel Models Using Stan.” *Journal of Statistical Software*, **80**(1), 1–28. doi:10.18637/jss.v080.i01.
- Couch S, Chang W, Gao C (2025). *mcptools: Model Context Protocol Servers and Clients for R*. URL <https://CRAN.R-project.org/package=mcptools>.
- Damnjanovic I, Reinschmidt K (2020). *Data Analytics for Engineering and Construction Project Risk Management*. Springer, Cham, Switzerland. doi:10.1007/978-3-030-14251-3.
- Elzhov TV, Mullen KM, Spiess AN, Bolker B (2016). *minpack.lm: R Interface to the Levenberg-Marquardt Nonlinear Least-Squares Algorithm Found in MINPACK, Plus Support for Bounds*. URL <https://CRAN.R-project.org/package=minpack.lm>.
- Fleming QW, Koppelman JM (2010). *Earned Value Project Management*. 4 edition. Project Management Institute, Newtown Square, PA.
- Frank J (2023). “PyCostTools: Python Tools for Cost Estimating and Learning Curves.” URL <https://github.com/frankij11/PyCostTools>.
- Gonçalves-Dosantos JC, García-Jurado I, Costa J (2022). *ProjectManagement: Management of Deterministic and Stochastic Projects*. R package, URL <https://CRAN.R-project.org/package=ProjectManagement>.
- Govan P (2014). *Causal Network Methods for Integrated Project Portfolio Risk Analysis*. Ph.D. thesis, Texas A&M University.
- Govan P, Damnjanovic I (2016). “The resource-based view on project risk management.” *Journal of Construction Engineering and Management*, **142**(9), 04016034. doi:10.1061/(ASCE)CE.1943-7862.0001157.
- Govan P, Damnjanovic I (2020). “Structural Network Measures for Risk Assessment of Construction Projects.” *ASCE-ASME Journal of Risk and Uncertainty in Engineering Systems, Part A: Civil Engineering*. doi:10.1061/AJRUA6.0001041.
- Højsgaard S (2012). “Graphical Independence Networks with the **gRain** Package for R.” *Journal of Statistical Software*, **46**(10), 1–26. doi:10.18637/jss.v046.i10.

- Iooss B, Da Veiga S, Janon A, Pujol G (2024). *sensitivity: Global Sensitivity Analysis of Model Outputs*. URL <https://CRAN.R-project.org/package=sensitivity>.
- Kelley D (2023). *plan: Tools for Project Planning*. R package, URL <https://CRAN.R-project.org/package=plan>.
- Luedeling E, Goehring L, Schiffrers K, Whitney C, Fernandez E (2023). *decisionSupport: Quantitative Support of Decision Making under Uncertainty*. R package, URL <https://CRAN.R-project.org/package=decisionSupport>.
- Ooms J (2014). “The jsonlite Package: A Practical and Consistent Mapping Between JSON Data and R Objects.” *arXiv:1403.2805 [stat.CO]*. URL <https://arxiv.org/abs/1403.2805>.
- Pearl J (1988). *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. Morgan Kaufmann, San Mateo, CA.
- Pearl J (2009). *Causality: Models, Reasoning, and Inference*. 2nd edition. Cambridge University Press, Cambridge, UK.
- Pouillot R, Delignette-Muller ML (2010). “Evaluating Variability and Uncertainty in Microbial Quantitative Risk Assessment Using Two R Packages.” *International Journal of Food Microbiology*, **142**(3), 330–340. doi:10.1016/j.ijfoodmicro.2010.07.011.
- Project Management Institute (2021). *A Guide to the Project Management Body of Knowledge (PMBOK Guide)*. 7th edition. Project Management Institute, Newtown Square, PA.
- R Core Team (2024). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.
- rOpenSci Statistical Software Peer Review (2021). *Statistical Software Peer Review: Standards and Guidelines*. URL <https://stats-devguide.ropensci.org>.
- Salvatier J, Wiecki TV, Fonnesbeck C (2016). “Probabilistic Programming in Python Using PyMC3.” *PeerJ Computer Science*, **2**, e55.
- Scutari M (2010). “Learning Bayesian Networks with the **bnlearn** R Package.” *Journal of Statistical Software*, **35**(3), 1–22. doi:10.18637/jss.v035.i03.
- Stan Development Team (2024). *RStan: The R Interface to Stan*. URL <https://mc-stan.org/>.
- Steward DV (1981). *Systems Analysis and Management: Structure, Strategy and Design*. Petrocelli Books, New York, NY.
- Torkia E (2023). “MCHammer.jl: Monte Carlo Simulation and Business Analysis Tools for Julia.” URL <https://github.com/etorkia/MCHammer.jl>.
- Vose D (2008). *Risk Analysis: A Quantitative Guide*. 3 edition. Wiley, Chichester.
- Wickham H, Cheng J, Jacobs A, Aden-Buie G, Schloerke B (2025). *ellmer: Chat with Large Language Models*. URL <https://ellmer.tidyverse.org>.

Affiliation:

Paul Govan
GE Aerospace
3290 Patterson Ave SE, Grand Rapids, MI 49512, USA
E-mail: paul.govan@geaerospace.com
URL: <https://paulgovan.com>